

Provenance-Preserving Evidence-Gated Knowledge-Graph Augmentation for Auditable Entity–Relation Extraction

Anonymous Authors

Abstract

Converting text into structured knowledge requires both accurate extraction and an explicit connection between predicted relations and their source evidence. Knowledge graphs can supply domain context, but a retrieved association is not necessarily a fact stated in the current sentence. This study proposes a provenance-preserving, evidence-gated framework that separates graph-informed candidate generation from output acceptance. A training-only graph provides bounded concept and relation hints to compact generative extractors, while deterministic span reconstruction, entity selection, and relation verification produce an evidence-linked assertion graph. Evaluation on ADE and CoNLL04 examines drug–adverse-effect extraction and general-domain entity–relation extraction within a shared character-span evaluation protocol. Dataset statistics characterize label concentration, sentence length, and split composition before the model comparisons. The completed test runs show higher relation F1 than the matched source-only generator on both benchmarks, with improved relation precision and explicit records of gate decisions. The findings demonstrate a practical mechanism for controlling graph augmentation and organizing source-grounded predictions for human review, while separating structural admissibility from semantic correctness.

Keywords: entity and relation extraction, knowledge graph, provenance, evidence grounding, adverse drug effects, selective prediction

1. Introduction

Structured information extraction connects unstructured text to searchable entities, typed relations, and downstream knowledge systems. In

4 applications related to safety, an extracted association must also be in-
5 spectable: users need to know which mention was identified, which sentence
6 supported the relation, and which transformations produced the final
7 assertion. Drug-related adverse-effect extraction illustrates this requirement
8 directly. A model that recognizes a drug and an adverse effect but invents
9 an association between them can create a misleading record even when
10 its output is syntactically valid. General-domain extraction provides a
11 complementary test of the same technical issue across a broader set of entity
12 and relation labels.

13 Neural entity and relation models provide effective approaches to the
14 recognition problem [1, 2]. Generative models offer a flexible alternative
15 in which a common structured output format can be reused across label
16 inventories. However, fluent generation does not guarantee that the predicted
17 objects are supported by their input [3]. This distinction matters when
18 extracted triples are inserted into a graph: errors become available to later
19 retrieval or reasoning stages, potentially propagating beyond the sentence
20 that generated them.

21 Knowledge graphs can contribute typed domain priors [4], and retrieval
22 augmentation can expose those priors without requiring all knowledge to be
23 stored in model parameters [5]. Their use introduces a specific authority
24 problem. A relation observed in a training sentence may suggest what to
25 search for, but it does not establish that the same relation holds in a new
26 sentence. A training graph may also return an example’s own annotation or
27 reveal content shared across data splits. The useful question is therefore not
28 only whether graph context increases recall, but whether its eligibility and
29 influence can be controlled and audited.

30 We address this problem with a staged architecture. Separate source-
31 only, exact-anchor, and high-recall graph-conditioned generators establish
32 a matched comparison of input conditions. A deterministic relation verifier
33 checks structural and source-evidence requirements. A complementary entity
34 gate uses agreement, source-matched anchors, or verified relation endpoints
35 to accept a subset of the graph-conditioned candidates. The final system
36 combines the accepted entities with verified relations whose endpoints sur-
37 vive. The architecture uses existing representation and adaptation methods;
38 its contribution is their governed composition, not a new backbone or a new
39 optimization algorithm.

40 This paper concentrates on ADE [6] and CoNLL04 [7]. ADE supplies a
41 drug-safety-related extraction task with two entity labels and one relation la-

bel. CoNLL04 supplies a general-domain control with four entity labels and five directed relation labels. Both use English sentence-level inputs, making it possible to examine the extraction and verification mechanisms without importing a long-document windowing problem. They are trained and evaluated separately; this is not a cross-dataset zero-shot transfer experiment.

The contributions are threefold:

- A provenance-aware graph-context policy that uses only training annotations, excludes candidates supported only by the current example, and clears graph hints for exact train/evaluation text overlaps.
- A deterministic composition of source-span reconstruction, entity selection, and relation verification that retains evidence-linked assertions and records the reasons behind output decisions.
- A reproducible two-benchmark analysis that combines split and label statistics, uniform strict-span scoring, validation ablations, completed test comparisons, and an exact-overlap sensitivity analysis.

The intended contribution is an auditable extraction layer for analyst-facing knowledge systems. It does not convert benchmark relation labels into clinical causality, establish calibrated risk probabilities, or validate autonomous safety decisions.

2. Related work

2.1. Joint entity and relation extraction

Entity recognition must resolve both mention boundaries and types; relation extraction additionally depends on recovering the correct directed endpoints [1]. Global consistency constraints have a long history in this setting, including the formulation associated with the CoNLL04 benchmark [7]. SpERT enumerates spans and jointly models entities and relations using a pretrained transformer [2]. It is a particularly relevant comparison because it performs the same joint task without requiring a generative JSON interface. Our experiments retrain SpERT from a clean base encoder on the same training partitions rather than using a released task checkpoint that may already include development data.

GLiNER and GLiREL provide label-conditioned alternatives for entity recognition and relation classification [8, 9]. These models differ from locally

75 fine-tuned generators in supervision, pretraining, and output construction.
 76 We therefore report the implemented train-calibrated pipeline as a named
 77 external baseline, not as an exhaustive assessment of either model family.
 78 A same-backbone zero-shot generator further separates the benefit of local
 79 adaptation from that of graph conditioning.

80 *2.2. Graph context and constrained generation*

81 Knowledge graphs represent typed entities and relations explicitly, while
 82 retrieval augmentation uses selected records as model context [4, 5]. GenIE
 83 constrains generative information extraction with a structured output space
 84 and entity/relation catalogues [10]. Catalogue validity and source support
 85 are different requirements: an admissible entity or relation may still not be
 86 asserted in the current input. Our graph context is consequently treated as
 87 a candidate prior whose eligibility depends on training provenance and local
 88 source matching.

89 BGE-M3 supplies frozen text representations for semantic pattern re-
 90 trieval [11]. It is neither a graph neural network nor a jointly optimized
 91 component of the extractor. QLoRA enables compact adaptation of the
 92 Qwen3 backbone [12, 13]. Neither method alone defines whether a candi-
 93 date should enter the final assertion graph; the proposed gates make that
 94 acceptance policy explicit.

95 *2.3. Evidence and safety-oriented knowledge construction*

96 Hallucination research distinguishes well-formed outputs from source-
 97 faithful ones [3]. In safety-related graph construction, prompt-based pro-
 98 cessing of accident reports likewise illustrates the need to trace generated
 99 knowledge back to its sources [14]. ADE offers a more controlled benchmark
 100 for a drug-safety-related relation task [6]. Nevertheless, an ADE label is
 101 an annotated textual association, not a patient-level causal conclusion. The
 102 present study evaluates extraction and implemented evidence constraints,
 103 rather than the medical or operational validity of downstream decisions.

104 **3. Datasets and distribution analysis**

105 *3.1. Benchmark sources and partitions*

106 We use the sentence-level ADE and CoNLL04 files distributed in SpERT
 107 format. Counts refer to that preprocessed distribution rather than every

108 sentence in the original source corpora. CoNLL04 retains its supplied train-
 109 ing, development, and test partitions. The ADE files contain 3,845 training
 110 sentences and 427 test sentences; a deterministic development slice of 384
 111 sentences is removed from the supplied training pool, leaving 3,461 train-
 112 ing sentences. The slice is selected by sorting SHA-256 hashes of the string
 113 formed from the fixed split seed `public-full-42` and each original record
 114 identifier. It is not selected by validation performance. This is one fixed ADE
 115 split, not an average over the multiple folds sometimes used in the literature.

116 Table 1 gives counts recomputed from the exact imported artifacts. The
 117 unit called a document in the software is one benchmark sentence, not a whole
 118 medical case report or news article. Both datasets are English; the results
 119 therefore support neither multilingual nor cross-lingual claims. Published
 120 annotations are imported rather than locally re-annotated.

Table 1: Dataset inventory for the exact experimental partitions. Length is the number of whitespace-delimited tokens in the reconstructed source text.

Dataset	Split	Sentences	Entities	Relations	Mean length
ADE	train	3461	8717	5470	21.17
ADE	validation	384	996	627	21.82
ADE	test	427	1126	724	21.24
CoNLL04	train	922	3377	1283	28.77
CoNLL04	validation	231	893	343	30.27
CoNLL04	test	288	1079	422	28.94

121 3.2. Entity and relation distributions

122 ADE contains `Drug` and `Adverse-Effect` entities and a single
 123 `Adverse-Effect` relation label. The imported endpoint convention is
 124 *adverse effect to drug*; we preserve it rather than reversing it to a more
 125 intuitive drug-to-effect direction. CoNLL04 contains `Loc`, `Org`, `Peop`,
 126 and `Other` entities, with `Work_For`, `Kill`, `OrgBased_In`, `Live_In`, and
 127 `Located_In` relations.

128 Figure 1 shows the training-label distributions, and [Appendix A](#) lists ex-
 129 act label counts for every partition. The single ADE relation label eliminates
 130 inter-class relation confusion but does not eliminate endpoint detection, di-
 131 rection, or false-pair errors. CoNLL04 introduces a broader relation inventory
 132 and more heterogeneous entity roles. Consequently, a higher aggregate F1

on ADE should not be interpreted as evidence that drug-safety reasoning is inherently easier than general-domain reasoning; the supervised extraction tasks differ in structure and size.




Figure 1: Training annotation distributions. Each panel uses its own entity or relation total as the denominator; entity and relation percentages are not pooled. Different colors and hatching distinguish the two benchmarks.

3.3. Sentence lengths and coverage

The mean training lengths are 21.17 reconstructed tokens for ADE and 28.77 for CoNLL04. Their test means are 21.24 and 28.94, respectively. The cumulative distributions in Figure 2 describe all three partitions rather than only the training set. These lengths are descriptive whitespace-token counts, not Qwen tokenizer counts. The largest observed reconstructed length is 97 for ADE and 118 for CoNLL04 across the imported partitions.

Every imported sentence contains at least one annotated relation. This is an important property of the selected benchmark files: they do not measure false-positive behavior over a deployment stream dominated by sentences with no relevant relations. We retain the supplied benchmark composition and do not create additional negative sentences for the reported tests.




Figure 2: Empirical cumulative sentence-length distributions by partition. Length is computed on deterministically reconstructed text; no label-informed window selection is used for these sentence-level inputs.

3.4. Split overlap and graph-context isolation

The graph is built separately for each dataset from its training annotations only. An exact-overlap check compares case-folded source strings after collapsing whitespace. ADE has no train/development or train/test exact-text overlap under this check. CoNLL04 has 14 development sentences and 19 test sentences whose text matches training content. For these sentences, both exact-anchor and high-recall graph context are cleared before generation.

Clearing graph hints does not undo exposure of the trained generator to duplicate training text. We therefore retain the supplied partitions for the main comparison and report an additional fixed-prediction test analysis that excludes the 19 overlapping CoNLL04 test sentences. We do not claim that the benchmark is duplicate-cluster isolated or that this exact-match check detects all near duplicates.

162 4. Proposed method

163 4.1. Source-grounded representation

164 Let a sentence x be the deterministic single-space concatenation of its
165 released tokens. An entity mention is represented as

$$e = (b, u, y), \quad x[b : u] = m, \quad (1)$$

166 where $[b, u]$ is a character interval, y is an entity label, and m is the copied
167 mention text. A relation $r = (e_s, z, e_t)$ has typed directed endpoints and
168 relation label z . Source identifiers and evidence spans are stored with these
169 objects. Repeated strings at different offsets remain distinct mentions. The
170 character-span requirement in Equation (1) prevents a normalized name
171 alone from receiving credit for the wrong occurrence.

172 The imported benchmark assigns every relation the status **explicit** and
173 uses its source sentence as the evidence span. This is a representation conven-
174 tion, not a new annotation of minimal causal evidence or uncertainty. Since
175 both endpoints occur in the same input sentence, evidence co-occurrence is
176 a relatively weak condition in these experiments.

177 4.2. Architecture and comparison systems

178 Figure 3 summarizes the implemented flow. SOE (Source-Only Extrac-
179 tion), EAE (Exact-Anchor Extraction), and HRGE (High-Recall Graph Ex-
180 traction) are separately trained QLoRA adapters on the same Qwen3-4B
181 snapshot. SOE uses source text and the training-derived ontology without
182 graph hints. EAE receives conservative concept hints, while HRGE receives
183 bounded anchors, edge priors, and semantic relation patterns.

184 EVGE (Evidence-Verified Graph Extraction) applies the relation verifier
185 to HRGE. CFE (Conservative-Fusion Extraction) pairs accepted HRGE en-
186 tities with endpoint-filtered raw HRGE relations. PGE (Provenance-Gated
187 Extraction) pairs the same accepted entities with endpoint-filtered verified
188 relations. These three systems are deterministic post-processing variants,
189 not additional trained networks. In deployment PGE needs EAE and HRGE
190 candidates; SOE is a comparison branch and does not participate in fusion.

191 4.3. Provenance-controlled graph context

192 Each training graph stores mention strings, types, directed relations, and
193 the sentence identifiers supporting them. For a training sentence d , a candi-
194 date is excluded if its only supporting provenance is d . Candidates with other




Figure 3: Two-benchmark architecture. Dataset-specific training graphs supply candidate context; the current sentence supplies evidence. Three independent adapters share the backbone configuration. Entity acceptance uses any of the listed signals; relation verification requires all checks. PGE retains only verified relations with surviving endpoints. Dashed boxes identify comparison paths, and no gate is a learned neural component.

supporting sentences remain eligible; their stored aggregate support statistics are not recomputed as a fully leave-one-out graph. This distinction prevents interpreting the provenance policy as complete removal of every contribution from the current example.

EAE uses at most 20 exact-matching concept hints with type balancing. HRGE uses three context channels. First, source-matched anchors must meet an empirical type-purity threshold of 0.8 and an English string-length threshold of four characters. Second, a typed edge prior requires support outside the current example and both endpoint strings in the current source. Third, frozen BGE-M3 retrieves semantically related relation-pattern records. If $\mathbf{H}$ contains normalized eligible record embeddings and $\mathbf{q}$ is the normalized source embedding, retrieval scores are

$$\mathbf{s} = \mathbf{H}\mathbf{q}, \quad \mathcal{I} = \text{TopK}_k\{i : s_i \geq \tau\}. \quad (2)$$

Equation (2) ranks candidate context, not accepted facts. The frozen HRGE limits are 12 anchors with at most two per type, six edge priors, and four semantic patterns with cosine similarity at least 0.72. The prompt explicitly identifies these objects as hints. Exact-overlap quarantine overrides retrieval and clears all graph context for the affected evaluation sentences.

4.4. Compact generation and deterministic reconstruction

The generator emits entity identifiers, copied text and types, plus relation identifiers, directed endpoint references, types, and explicit status. It is not asked to calculate source offsets. Deterministic expansion assigns copied mentions to occurrences in the supplied sentence and reconstructs character spans; unmatched strings and invalid references are recorded. This separates language-model candidate generation from source-coordinate bookkeeping.

All three generators use one training epoch, batch size one, gradient accumulation four, and learning rate 2×10^{-4} . QLoRA uses rank eight, alpha 16, dropout 0.05, NF4 double quantization, and bfloat16 computation. Attention query/key/value/output projections and feed-forward gate/up/down projections are adapted; the backbone remains frozen. AdamW uses the PyTorch defaults for its other parameters, without a learning-rate scheduler or warm-up. The maximum training sequence is 4,096 tokens, and prompt tokens are masked from the loss. Training uses the same dataset partition for all three branches and seed 42 throughout the primary comparisons.

Inference uses the recorded 4,096-token input limit, 1,024-token generation budget, and 90-second job limit, with two inference workers on the local

GPU. The backbone revision and all input hashes are recorded in the artifact manifest. Failed jobs are represented by empty predictions, not removed from the denominator. These experiments use the full supplied sentence, without annotation-guided rescue windows or label-dependent sentence truncation.

4.5. Relation verification and entity acceptance

For each relation label z , training annotations define permitted source-type and target-type sets S_z and T_z . The implemented type constraint is $y_s \in S_z$ and $y_t \in T_z$; it does not additionally require that every specific source/target type pair has been observed together. The verifier $V(r)$ also requires existing distinct entity references, a known relation label, explicit status, valid evidence, and both endpoint strings in a common source-evidence span. Its accepted set is

$$R_V = \{r \in R_H : V(r) = 1\}. \quad (3)$$

The entity gate uses three Boolean signals:

$$A(e) = A_{\text{agree}}(e) \vee A_{\text{anchor}}(e) \vee A_{\text{endpoint}}(e), \quad E_A = \{e \in E_H : A(e) = 1\}. \quad (4)$$

Agreement means that EAE predicts the same case-folded, whitespace-normalized text and type. Anchor support requires a same-type source-gated HRGE anchor. Endpoint support means that the entity participates in a relation retained by Equation (3). The three signals in Equation (4) are complementary but not statistically independent; verified endpoints are a rule-based signal, not a second human judgement. Duplicate normalized text/type candidates are collapsed within the sentence.

Let $F(R, E)$ retain relations in R only when both referenced entities survive in E . The two compositions are

$$G_{\text{CFE}} = (E_A, F(R_H, E_A)), \quad G_{\text{PGE}} = (E_A, F(R_V, E_A)). \quad (5)$$

Equation (5) makes the difference between raw-relation fusion and verified fusion explicit. Accept/reject decisions are logged. A retained relation satisfies the stated rules, but the rules alone cannot establish whether the sentence asserts the proposed meaning.

256 5. Experimental protocol

257 5.1. Research questions and scope

258 We ask whether PGE improves strict entity and relation extraction over
259 SOE; how graph context and deterministic selection contribute on develop-
260 ment data; and whether observed gains persist after exact-overlap sentences
261 are removed. External baselines establish whether the generative design is
262 competitive with other extraction families, rather than assuming that im-
263 provements over SOE imply state-of-the-art performance.

264 The two-benchmark reporting scope was narrowed retrospectively after
265 validation and available test results had been inspected. It is a focused anal-
266 ysis, not a preregistered or representative sample of all attempted extraction
267 domains. Within the existing protocol, model and gate settings were frozen
268 before test release. SOE and PGE were promoted when validation relation
269 F1 did not regress by more than one percentage point and the required com-
270 pletion and evidence checks passed. This was a pragmatic point-estimate
271 promotion rule, not a statistical non-inferiority test. The current paper uses
272 the completed seed-42 tests and does not infer multi-seed stability from them.

273 5.2. Baselines and evaluation parity

274 The internal development comparison includes all six named systems.
275 The test table includes the two promoted systems and three completed ex-
276 ternal comparisons. Qwen3-4B zero-shot uses the same backbone without lo-
277 cal adapters, followed by the recorded source expansion and relation verifier.
278 SpERT is trained from `bert-base-cased` on the same training partitions
279 for 20 epochs, with batch size two and a relation threshold of 0.4. Its final
280 checkpoint is used without selecting a test-best epoch. GLiNER+GLiREL
281 uses predicted, not gold, entities and relation calibration confined to an inner
282 training subset; the entity threshold is 0.5. These are heterogeneous baseline
283 configurations, not equal-compute experiments.

284 The recorded GLiREL calibration selects canonical relation labels and a
285 relation threshold of 0.0 for ADE, and naturalized labels with threshold 0.1
286 for CoNLL04. No development or test gold is used to choose those settings.

287 All reported main-table predictions are rescored with one public
288 character-span evaluator. An entity key contains its exact source interval
289 and type. A relation key contains both typed directed entity keys and the
290 relation label. Predicted evidence offsets must resolve to the supplied source;
291 the evaluator does not recover missing offsets by searching normalized

292 text. Duplicate keys receive one vote, invalid objects remain unmatched
 293 predictions, and omitted jobs remain empty. For pooled counts,

$$P = \frac{TP}{TP + FP}, \quad R = \frac{TP}{TP + FN}, \quad F_1 = \frac{2TP}{2TP + FP + FN}. \quad (6)$$

294 Equation (6) is applied across sentences within each dataset, not after av-
 295 eraging their individual F1 scores. Because all imported statuses are ex-
 296 plicit, claim-status-aware scores do not form a separate uncertainty bench-
 297 mark. The unified evaluator removes small denominator differences caused
 298 by legacy duplicate handling; unchanged source artifacts and a reconciliation
 299 report preserve the earlier metrics.

300 5.3. Statistical and compliance analyses

301 We compute 20,000 paired bootstrap resamples of test sentences, using
 302 seed 20260830 and fixed predictions, and recompute pooled F1 within each
 303 resample. The percentile interval describes sentence-sampling uncertainty for
 304 the completed model pair, not variability across training seeds. Intervals are
 305 descriptive, unadjusted for multiple comparisons, and are not used to change
 306 the frozen models or gates.

307 Evidence coverage and rule-invalid counts are measured separately from
 308 extraction F1. These compliance counts refer to emitted prediction objects
 309 before the evaluator’s key deduplication. Exact-overlap sensitivity removes
 310 the same affected test sentence identifiers from both systems, without retrain-
 311 ing or regenerating predictions. Test labels are used for reporting statistics
 312 and scoring only, not for constructing graph hints or choosing thresholds in
 313 this revision.

314 6. Results

315 6.1. Completed test comparisons

316 Table 2 reports the completed test results under the common evaluator.
 317 PGE raises ADE relation precision from 80.62% to 87.00% and recall from
 318 67.82% to 72.10%, giving relation F1 of 78.85% compared with 73.67% for
 319 SOE. Entity F1 remains essentially unchanged, at 87.93% versus 87.97%.
 320 Thus, the main improvement is relational rather than a general gain in men-
 321 tion recognition.

322 On CoNLL04, PGE raises relation precision from 54.61% to 62.40% and
 323 recall from 56.16% to 57.82%. Relation F1 increases from 55.37% to 60.02%,

324 while entity F1 decreases from 83.83% to 83.20%. The composition trades
 325 some entity coverage for a better relation precision–recall balance in this run.

Table 2: Completed test comparisons, strict source-character-span micro scores (%). Qwen3-ZS is Qwen3-4B zero-shot with the recorded verifier; GLiNER+GLiREL uses training-only calibration. Each dataset uses its own test partition. Locally trained systems use seed 42; all rows are rescored with the same evaluator.

Dataset	System	Ent. P	Ent. R	Ent. F1	Rel. P	Rel. R	Rel. F1
ADE	Qwen3-ZS + verifier	63.83	63.32	63.58	39.36	22.24	28.42
ADE	GLiNER + GLiREL	79.30	64.65	71.23	60.31	43.23	50.36
ADE	SpERT	90.65	92.10	91.37	82.19	86.05	84.08
ADE	SOE	92.21	84.10	87.97	80.62	67.82	73.67
ADE	PGE	92.89	83.48	87.93	87.00	72.10	78.85
CoNLL04	Qwen3-ZS + verifier	48.78	53.94	51.23	30.92	19.19	23.68
CoNLL04	GLiNER + GLiREL	60.78	14.37	23.24	13.79	3.79	5.95
CoNLL04	SpERT	88.61	89.43	89.02	71.01	68.48	69.72
CoNLL04	SOE	85.83	81.93	83.83	54.61	56.16	55.37
CoNLL04	PGE	86.04	80.54	83.20	62.40	57.82	60.02

326 SpERT has the highest entity and relation F1 among these completed
 327 systems: 91.37%/84.08% on ADE and 89.02%/69.72% on CoNLL04. PGE
 328 therefore improves the matched generative baseline but does not surpass this
 329 trained span-based model. PGE also exceeds the recorded zero-shot and
 330 GLiNER+GLiREL pipelines on both extraction tasks. The latter observa-
 331 tion is specific to their implemented label mapping and calibration; it is not a
 332 claim that every configuration of those model families would perform worse.

333 6.2. Paired differences and uncertainty

334 Using unrounded counts, PGE-minus-SOE relation F1 is 5.18 percentage
 335 points on ADE, with a paired 95% interval of [0.31, 10.29], and 4.65 points
 336 on CoNLL04, with an interval of [0.04, 9.12]. The positive point estimates
 337 are consistent with the test-table improvements, but the intervals are broad
 338 and their lower bounds are close to zero. We interpret them as evidence
 339 about these completed test predictions, not as a general guarantee across
 340 seeds. Entity-F1 differences are -0.04 points on ADE, interval $[-1.75, 1.62]$,
 341 and -0.63 on CoNLL04, interval $[-2.38, 1.04]$; neither supports an entity-F1
 342 improvement claim.

343 6.3. Development-set component analysis

344 Table 3 separates candidate generation from deterministic selection on
 345 development data. On ADE, HRGE has relation F1 of 76.29%, above SOE’s

73.16%. EVGE is identical to HRGE, and PGE is identical to CFE, showing that the relation verifier removes no additional relations in these variants. The entity gate reduces HRGE recall and gives PGE relation F1 of 75.91%. It is therefore incorrect to attribute all ADE gains to verification or to claim that every stage independently improves F1.

On CoNLL04, HRGE relation F1 is 58.26%, below SOE’s 60.24%. Verification raises it to 59.60%, and the final composition reaches 60.19%. This is effectively unchanged from SOE at the development point estimate, rather than a clear development-set gain. Entity precision rises after gating, while recall and F1 decrease. The distinction between development and test behavior is retained rather than replacing the ablation with a test-selected variant.

Table 3: Development-set ablation under the common strict-span evaluator (micro %, seed 42). These rows characterize the frozen architecture and are not additional promoted test systems.

Dataset	System	Ent. P	Ent. R	Ent. F1	Rel. P	Rel. R	Rel. F1
ADE	SOE	89.25	82.53	85.76	78.82	68.26	73.16
ADE	EAE	90.49	81.22	85.61	80.84	67.30	73.46
ADE	HRGE	90.35	85.54	87.88	80.04	72.89	76.29
ADE	EVGE	90.35	85.54	87.88	80.04	72.89	76.29
ADE	CFE	90.48	83.94	87.08	81.39	71.13	75.91
ADE	PGE	90.48	83.94	87.08	81.39	71.13	75.91
CoNLL04	SOE	84.94	82.08	83.49	61.33	59.18	60.24
CoNLL04	EAE	83.71	82.87	83.29	58.10	60.64	59.34
CoNLL04	HRGE	82.35	81.52	81.94	60.06	56.56	58.26
CoNLL04	EVGE	82.35	81.52	81.94	62.99	56.56	59.60
CoNLL04	CFE	84.97	78.50	81.61	63.37	55.98	59.44
CoNLL04	PGE	84.97	78.50	81.61	65.08	55.98	60.19

6.4. Rule compliance and its interpretation

The test SOE output contains 31 rule-invalid relation objects among 610 ADE objects and 16 among 436 CoNLL04 objects. PGE retains 600 and 392 relation objects, respectively, with zero rule-invalid objects. HRGE already has zero invalid ADE relations and only three invalid CoNLL04 relations among 413 objects. These counts distinguish the contribution of candidate quality from that of the verifier.

The measured unsupported-evidence count is zero for SOE, HRGE, and PGE on both test datasets. Because each relation’s evidence is the whole sentence, co-occurrence is easy to satisfy once endpoints resolve. Thus, these experiments do not demonstrate a reduction from a nonzero unsupported-claim rate. They show explicit evidence binding and rule compliance, while

the gold-referenced F1 scores measure a different property: whether the predicted typed relation agrees with the annotation.

6.5. *Exact-overlap sensitivity*

Removing the 19 exact train/test overlaps leaves 269 CoNLL04 test sentences, 1,023 gold entities, and 403 gold relations. SOE relation F1 becomes 53.79% and PGE becomes 59.10%, a 5.31-point difference; entity F1 is 83.02% and 82.90%, respectively. The observed relation advantage therefore persists without those exact overlaps. This analysis addresses exact repeated text, not all forms of semantic overlap or pretraining contamination. ADE has no such overlaps, so its sensitivity result is identical to the full test.

7. Discussion

7.1. *What the two benchmarks establish*

The common result is a relation-F1 improvement over source-only generation, not universal superiority over extraction architectures. The gains arise alongside higher relation precision on both tests, while entity F1 stays similar or declines slightly. This is consistent with an architecture designed to govern candidate acceptance rather than simply maximize output volume. The validation ablation further shows that graph conditioning and selection interact: recall-oriented context can help on ADE but does not automatically improve CoNLL04, and a verifier may be inactive when candidates already satisfy its rules.

The distribution analysis helps interpret these differences. ADE has a larger training pool, shorter average sentences, and a single relation class. CoNLL04 has fewer training sentences, longer inputs, and five relation labels. These factors are confounded rather than experimentally isolated. Their association with the observed scores motivates separate dataset reporting; it does not establish a causal explanation for the performance gap.

7.2. *Auditability as a distinct design objective*

A system can be less accurate than SpERT while still exposing a useful acceptance mechanism for generated graph objects. The contribution is strongest when the entire path is inspectable: graph support identifies why a candidate was suggested; reconstructed offsets identify where it appears; gate records identify why it was retained or rejected. However, this benefit is architectural. No user study in the present experiments measures reduced

analyst effort, improved decisions, or greater trust. It would be inappropriate to substitute perfect rule compliance for such evidence.

The gates also impose a coverage cost. Agreement operates on normalized text/type pairs, and repeated mentions may be collapsed during fusion even though the evaluator distinguishes offsets. This can remove legitimate occurrences. A legal endpoint signature and sentence-level co-occurrence are insufficient to prove a relation’s meaning, and three gate signals should not be interpreted as three independent confirmations. These observations explain why the final output remains a candidate assertion graph for review.

7.3. Implications for safety-related applications

ADE links the experiments to drug-safety information extraction, while CoNLL04 tests general relational structure. The framework can present an analyst with the candidate relation, source sentence, typed endpoints, and provenance of its acceptance path. An analyst can then accept, amend, or reject the proposed association. Absence from the retained graph is not evidence of the absence of an adverse effect or another real-world relation.

The reported task stops before pharmacovigilance judgement, patient-specific causal attribution, or operational action. Neither dataset validates risk propagation, resilience assessment, or autonomous intervention. Extending this work toward those applications requires task-specific outcome measures and expert assessment in addition to extraction F1.

8. Limitations and responsible use

The focused, retrospectively narrowed dataset scope and single completed training seed limit generalization. The fixed ADE partition is not directly interchangeable with published multi-fold results. Both corpora are English and consist of relation-bearing sentences; multilingual behavior, long-report processing, and false-positive rates on unrestricted text streams are outside the evaluated setting. Sentence-level bootstrap intervals cannot recover dependence between sentences originating from the same upstream article when those article identifiers are unavailable.

Exact-overlap quarantine prevents direct graph hints for known repeated source text but does not remove the generator’s training exposure. The sensitivity analysis excludes known exact duplicates only. Frozen pretrained encoders and generators may also have seen related material during pretraining, which the available artifacts cannot rule out. The ontology’s allowed source

438 and target sets are less restrictive than a catalogue of observed joint type
439 pairs.

440 The one-epoch generative systems, 20-epoch SpERT, and pretrained
441 label-conditioned models do not form a compute-matched benchmark. PGE
442 requires two adapted generators plus frozen retrieval, so compact adaptation
443 must not be confused with measured end-to-end cost savings. No new
444 clinical annotations or independent evidence-judgement labels were collected
445 for this paper. Public availability of a dataset or model does not by itself
446 permit every downstream redistribution or commercial use; original licenses
447 remain applicable.

448 9. Conclusions

449 This study presents a provenance-preserving, evidence-gated graph-
450 augmentation framework and evaluates it on ADE and CoNLL04. Dataset
451 statistics expose the label and input-length differences underlying the two
452 benchmarks. Under a unified strict-span evaluator, the completed PGE
453 tests improve relation F1 over SOE by 5.18 and 4.65 percentage points,
454 respectively, while offering an explicit source-linked acceptance path. The
455 CoNLL04 relation advantage remains after exact train/test overlaps are
456 excluded. SpERT remains stronger in overall extraction F1, and the
457 development ablations show that the gates are not uniformly beneficial in
458 isolation. The supported outcome is controlled, auditable graph-informed
459 extraction for human review, rather than a claim of universal accuracy
460 dominance or validated autonomous safety reasoning.

461 Acknowledgements and funding

462 Author-identifying funding and acknowledgement information is supplied
463 on the separate title page for double-blind review.

464 Declaration of competing interest

465 The authors declare that they have no known competing financial in-
466 terests or personal relationships that could have appeared to influence this
467 work.

468 Data and code availability

469 The experiments use the public ADE and CoNLL04 distributions in
 470 SpERT format. An anonymized reproducibility package will contain the
 471 split construction, dataset statistics, source hashes, evaluation code, and
 472 derived result tables. Dataset text and model weights remain subject to
 473 their original licenses. The identifying repository URL is withheld during
 474 double-blind review.

475 Declaration of generative AI and AI-assisted technologies in the 476 writing process

477 The authors used OpenAI Codex to assist with code inspection, language
 478 editing, figure preparation, and manuscript organization. The authors are
 479 responsible for reviewing the resulting material and for the final content.
 480 The language models evaluated as research systems are described separately
 481 in the Methods.

482 Appendix A. Label counts by partition

483 Table A.4 provides the counts underlying the distribution analysis. En-
 484 tity and relation labels with the same string in ADE belong to different
 annotation inventories and must not be combined.

Table A.4: Exact entity and relation annotation counts in the imported files.

Dataset	Kind	Label	Train	Validation	Test
ADE	entity	Adverse-Effect	4631	529	616
ADE	entity	Drug	4086	467	510
ADE	relation	Adverse-Effect	5470	627	724
CoNLL04	entity	Loc	1219	322	427
CoNLL04	entity	Org	616	170	198
CoNLL04	entity	Other	455	118	133
CoNLL04	entity	Peop	1087	283	321
CoNLL04	relation	Kill	179	42	47
CoNLL04	relation	Live_In	330	91	100
CoNLL04	relation	Located_In	247	65	94
CoNLL04	relation	OrgBased_In	271	76	105
CoNLL04	relation	Work_For	256	69	76

485

486 References

- 487 [1] J. Li, A. Sun, J. Han, C. Li, A survey on deep learning for named entity
488 recognition, *IEEE Transactions on Knowledge and Data Engineering*
489 34 (1) (2022) 50–70. doi:[10.1109/TKDE.2020.2981314](https://doi.org/10.1109/TKDE.2020.2981314).
- 490 [2] M. Eberts, A. Ulges, Span-based joint entity and relation extraction
491 with transformer pre-training, in: *ECAI 2020*, IOS Press, 2020. doi:
492 [10.3233/FAIA200321](https://doi.org/10.3233/FAIA200321).
- 493 [3] Z. Ji, N. Lee, R. Frieske, T. Yu, D. Su, Y. Xu, E. Ishii, Y. J.
494 Bang, A. Madotto, P. Fung, Survey of hallucination in natural lan-
495 guage generation, *ACM Computing Surveys* 55 (12) (2023) 1–38. doi:
496 [10.1145/3571730](https://doi.org/10.1145/3571730).
- 497 [4] A. Hogan, E. Blomqvist, M. Cochez, C. D’Amato, G. de Melo, C. Gutier-
498 rez, S. Kirrane, J. E. L. Gayo, R. Navigli, S. Neumaier, A.-C. N.
499 Ngomo, A. Polleres, S. M. Rashid, A. Rula, L. Schmelzeisen, J. Se-
500 queda, S. Staab, A. Zimmermann, Knowledge graphs, *ACM Computing*
501 *Surveys* 54 (4) (2021) 1–37. doi:[10.1145/3447772](https://doi.org/10.1145/3447772).
- 502 [5] P. Lewis, E. Perez, A. Piktus, F. Petroni, V. Karpukhin, N. Goyal,
503 H. Küttler, M. Lewis, W. tau Yih, T. Rocktäschel, S. Riedel, D. Kiela,
504 [Retrieval-augmented generation for knowledge-intensive NLP tasks](#), in:
505 *Advances in Neural Information Processing Systems*, Vol. 33, 2020, pp.
506 9459–9474.
507 URL [https://proceedings.neurips.cc/paper/2020/hash/6b49323](https://proceedings.neurips.cc/paper/2020/hash/6b493230205f780e1bc26945df7481e5-Abstract.html)
508 [0205f780e1bc26945df7481e5-Abstract.html](https://proceedings.neurips.cc/paper/2020/hash/6b493230205f780e1bc26945df7481e5-Abstract.html)
- 509 [6] H. Gurulingappa, A. M. Rajput, A. Roberts, J. Fluck, M. Hofmann-
510 Apitius, L. Toldo, Development of a benchmark corpus to support the
511 automatic extraction of drug-related adverse effects from medical case
512 reports, *Journal of Biomedical Informatics* 45 (5) (2012) 885–892. doi:
513 [10.1016/j.jbi.2012.04.008](https://doi.org/10.1016/j.jbi.2012.04.008).
- 514 [7] D. Roth, W. tau Yih, [A linear programming formulation for global infer-](#)
515 [ence in natural language tasks](#), in: *Proceedings of the Eighth Conference*
516 *on Computational Natural Language Learning*, Association for Compu-
517 tational Linguistics, 2004, pp. 1–8.
518 URL <https://aclanthology.org/W04-2401/>

- [8] U. Zaratiana, N. Tomeh, P. Holat, T. Charnois, GLiNER: Generalist model for named entity recognition using bidirectional transformer, in: Proceedings of the 2024 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Association for Computational Linguistics, 2024, pp. 5364–5376. doi:10.18653/v1/2024.naacl-long.300.
- [9] J. Boylan, C. Hokamp, D. G. Ghalandari, GLiREL: Generalist model for zero-shot relation extraction, arXiv preprint arXiv:2501.03172 (2025). doi:10.48550/arXiv.2501.03172.
- [10] M. Josifoski, N. D. Cao, M. Peyrard, F. Petroni, R. West, [GenIE: Generative information extraction](#), in: Proceedings of the 2022 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Association for Computational Linguistics, 2022, pp. 4626–4643. doi:10.18653/v1/2022.naacl-main.342.
URL <https://aclanthology.org/2022.naacl-main.342/>
- [11] J. Chen, S. Xiao, P. Zhang, K. Luo, D. Lian, Z. Liu, BGE M3-embedding: Multi-lingual, multi-functionality, multi-granularity text embeddings through self-knowledge distillation, arXiv preprint arXiv:2402.03216 (2024). doi:10.48550/arXiv.2402.03216.
- [12] T. Dettmers, A. Pagnoni, A. Holtzman, L. Zettlemoyer, [QLoRA: Efficient finetuning of quantized LLMs](#), in: Advances in Neural Information Processing Systems, Vol. 36, 2023, pp. 10088–10115.
URL https://proceedings.neurips.cc/paper_files/paper/2023/hash/1feb87871436031bdc0f2beaa62a049b-Abstract-Conference.html
- [13] Qwen Team, Qwen3 technical report, arXiv preprint arXiv:2505.09388 (2025). doi:10.48550/arXiv.2505.09388.
- [14] Q. Liu, X. Luo, [Large language model-based knowledge graph construction of unstructured accident reports to improve construction safety](#), in: EG-ICE 2025: AI-Driven Collaboration for Sustainable and Resilient Built Environments, University of Strathclyde Publishing, 2025, pp. 536–546. doi:10.17868/strath.00093269.
URL <https://strathprints.strath.ac.uk/93269/>